Smart Grid Energy Monitoring & Billing

A Lambda Architecture Data Platform

G.W. Niranga Niluminda (Dev A) Dev B

EC8203 Applied Big Data Engineering — Mini Project Report

Contents

1. Introduction 1
2. Architecture Decision: Lambda vs Kappa 2
3. Designed Architecture 4
4. Technology Stack 5
5. Data Ingestion Implementation 7
6. Processing Layer Implementation 8
7. Storage & Serving Layer 10
8. Observability Design 10
9. Results 11
10. Limitations, Trade-offs, and Production-Scale Changes 12
Appendix 14

1. Introduction

1.1 Use case

This project implements **Use Case 3 — Smart Grid Energy Monitoring & Billing** from the assignment brief. A utility company wants two things at once:

1. **Live visibility** into grid load and the renewable (solar) contribution to that load, broken down by grid zone, updated continuously as smart-meter readings arrive.
2. A **daily, authoritative reconciliation**: once a business day has closed and the day's tariff and weather reference data has been received from the billing system, compute exactly what each household owes and how much of the zone's demand was met by solar that day.

Business question. *What is the current grid load and renewable contribution by zone, and what will each household's bill look like once daily tariff data is applied to their consumption?*

These are genuinely different problems even though they share a data source. The first needs to be fast and can tolerate being slightly wrong (a reading that arrives a few seconds late does not materially change "current load"). The second needs to be exact, auditable, and reproducible — it is a financial calculation that must reconcile to the cent and must be re-computable if a tariff is corrected after the fact. That distinction is the entire justification for the architecture chosen in §2.

1.2 Data sources (simulated)

Two Python services simulate the real-world sources described in the brief:

Source	Frequency	Fields
Smart meters (`meter-simulator`)	one reading per household roughly every 2 simulated seconds	`meter_id, household_id, grid_zone, power_consumption_kwh, solar_generation_kwh, reading_interval_seconds, event_timestamp`
Tariff & weather feed (`batch-file-generator`)	once per simulated day	`tariff_reference_<date>.csv` (household_id, tariff_rate, billing_tier, subsidy_flag, effective_date) and `weather_forecast_<date>.csv` (grid_zone, forecast_date, cloud_cover_pct, temp_c, solar_irradiance_index)

The streaming source deliberately injects realistic anomalies at a configurable rate (`SIM_ANOMALY_RATE`, default 2%): dropped readings, negative values, implausible spikes, a "stuck" meter repeating its last value, and zero readings. The platform is expected to detect and quarantine these, not silently average them into the results (see §5 and §6).

1.3 Assumptions and simplifications

- **Simulated time.** One simulated business day is compressed into `SIM_DAY_SECONDS` wall-clock seconds (600 by default — 10 minutes; the demo recording uses a shorter value so a full day-close is visible inside a short video). All services derive "what day is it" from a single shared `SimClock`, so the streaming source, the daily feed, and the batch orchestration all agree on the current simulated date.
- **Household population.** 200 synthetic households spread evenly across 6 grid zones (`data/seed/households.csv`), each with a diurnal load profile and (for ~35% of households) a solar installation, deterministically generated from a seed so results are reproducible.
- **Currency and tariff model.** Tariffs are illustrative (LKR/kWh), with three billing tiers (flat, tiered-block, subsidised) chosen per household at random each simulated day, to exercise the Strategy-pattern billing logic described in §6.
- **Single-node deployment.** Kafka runs as a single KRaft broker, Spark runs in local mode inside each processing container, and Airflow runs in `standalone` mode (webserver + scheduler + triggerer in one process). This is the right scope for a mini-project demonstrating the architecture; §10 discusses what would change at production scale.
- **No authentication** on the serving API or dashboards, consistent with a local/demo deployment.

---

2. Architecture Decision: Lambda vs Kappa

(Full written decision record: see ADR-0001 in `docs/decisions/`*.)*

2.1 The forces

	Real-time zone view	Daily household bill
Latency requirement	seconds	"ready the morning after"
Consistency requirement	eventual / approximate acceptable	strong — must reconcile exactly
Input	unbounded, partial stream	one bounded, closed day + a once-daily reference feed
Reprocessing	not meaningful	**required** (a corrected tariff or late data must be re-billable)
Natural trigger	continuous	"business day D has closed" **and** "the feed for D has landed"
Computation	rough per-zone sums over a short window	a per-household join of a full day's consumption × tariff × subsidy, minus a solar export credit

These are not the same computation at two different freshness levels — they are two different computations with two different correctness models and two different consumers. That is precisely the situation Lambda architecture was designed for.

2.2 Decision: Lambda

We adopted a **Lambda architecture**:

- A **speed layer** (Apache Spark Structured Streaming) consumes the raw meter-reading stream and maintains an approximate, continuously-updated view of load and renewable share per zone, using event-time windowing with watermarking to bound how long it waits for late data.
- A **batch layer** (Apache Spark batch job, orchestrated by Apache Airflow) runs once a simulated day's readings and its tariff/weather feed have both landed. It re-derives the authoritative per-household bill and per-zone solar summary from the **immutable master dataset** — every raw reading ever ingested, partitioned by business date in Parquet — so that a corrected tariff or a late-arriving batch of readings can always be reconciled by simply re-running the job for that date.
- A **serving layer** (PostgreSQL + a read-only FastAPI service) merges the two: a day that has not yet been through the batch layer is served as *provisional* (speed-layer numbers); a day the batch layer has finalised is served as *final* (batch-layer numbers).

2.3 Rejected alternative: Kappa

A Kappa architecture would treat the daily tariff feed as a compacted Kafka topic and maintain the household bill as continuously-updated streaming state, using the retained log for reprocessing instead of a separate batch job. We considered this and rejected it for this use case, for three reasons:

1. **Billing is inherently periodic and bounded**, not continuous. A bill is "the invoice for day D," not a number that is always slightly provisional. Expressing "the day has closed, emit the final invoice" as watermark/allowed-lateness control flow inside a single streaming job is harder to reason about, and harder to audit, than a scheduled DAG with an explicit `business_date` parameter, a recorded run, and retries.
2. **Reprocessing cost.** Correcting one historical tariff error under Kappa means replaying the *entire* meter-reading history through a stateful streaming job. Under Lambda, the

same correction touches exactly one Parquet partition — the closed day in question — and nothing else.

3. **The assignment's learning objectives and preferred stack** explicitly ask for both stream *and* batch processing techniques, and name Airflow for "managing batch jobs or reporting pipelines." Kappa would collapse the exercise into a single paradigm and leave the batch/orchestration objectives undemonstrated.

Kappa's usual advantage — one codebase, no dual logic to keep in sync — is weaker here than it looks, because the two layers are not doing the same computation twice: the speed layer computes a rough, live aggregate; the batch layer computes an exact, audited one from different inputs (it also needs the tariff feed, which does not exist until the day closes). To prevent the remaining shared logic (reading validation, the energy-balance split between self-consumption and grid import) from drifting between the two layers, both the speed and batch implementations import the *same* pure functions from a shared library (`energy_common.domain.transforms`) rather than re-implementing them — see §6.

When would we choose Kappa instead? If the requirement were only the real-time zone dashboard, with no daily authoritative bill, Kappa's simplicity would win outright. The daily bill's need for periodic, bounded, auditable recomputation is what tips the decision to Lambda.

---

3. Designed Architecture

(Source diagrams: `docs/ARCHITECTURE.md`.)

3.1 Context

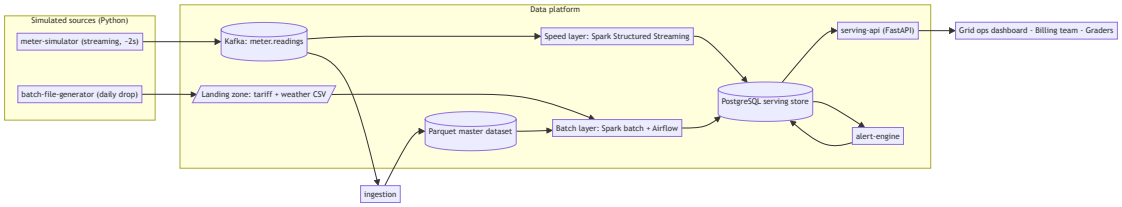


Figure 1: System context: simulated sources, the data platform, and consumers.

3.2 The two layers and the serving merge

The speed layer validates and enriches each reading, then computes tumbling 1-minute and sliding 5-minute/1-minute windowed aggregates per zone (`total_consumption_kwh`, `total_solar_kwh`, `renewable_pct`, `active_meter_count`), upserted into `rt_zone_metrics`.

The batch layer reads one full business day's partition from the Parquet master dataset, rolls it up per household, joins it against that day's tariff and weather reference files, applies a billing strategy (flat / tiered / subsidised — see §6), nets off a solar export credit, and writes `daily_household_billing` and `daily_zone_solar_summary`. It also renders a human-readable HTML/CSV report artifact.

`serving-api` never computes; it only reads. For "today," it returns the speed-layer numbers tagged `"status": "provisional"`. For any day the batch layer has completed (recorded in a `batch_runs` audit table), it returns the batch-layer numbers tagged `"status": "final"`. This tagging is the concrete, API-visible expression of the Lambda decision in §2 — a client can always tell which layer answered its request.

3.3 One simulated day, sequenced

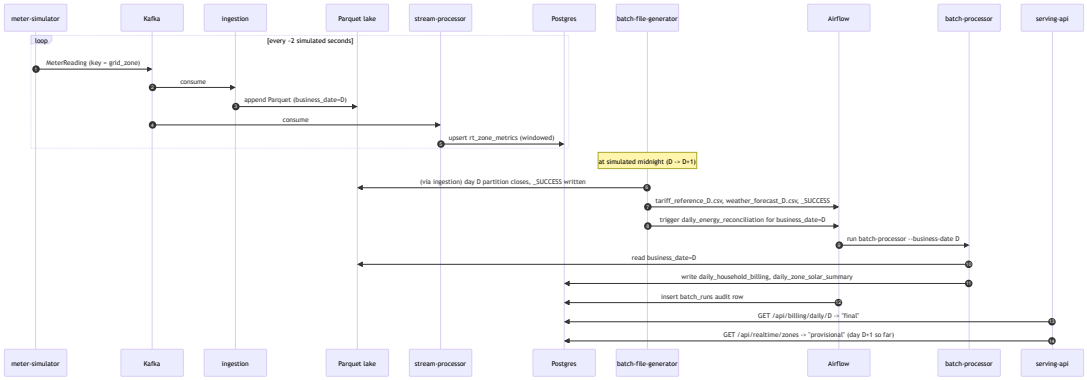


Figure 2: One simulated day, sequenced: streaming ingestion through to the daily batch reconciliation and serving.

3.4 Hexagonal service structure

Every service (both Part 1 and Part 2) is organised the same way, so understanding one means understanding all of them:

```
services/<name>/src/<pkg>/
  domain/        pure business rules + models + port interfaces (no framework
↪  imports)
  application/   use-cases that orchestrate the domain via ports
  adapters/      inbound/outbound implementations of those ports (Kafka,
↪  Postgres, HTTP, Spark)
  config.py      one Settings class per service, env-driven, validated at
↪  startup
  composition.py the one place that wires concrete adapters into a use-case
  __main__.py    process entrypoint
```

The domain layer defines interfaces (`Protocol` classes); the adapters implement them; only the composition root imports both sides. This is what makes the business logic — window aggregation, billing calculation, alert rules — unit-testable with plain fakes, with zero Kafka/Postgres/Spark running.

3.5 Deployment

`make up` starts Part 1 only; `make up-full` layers Part 2 on top via Docker Compose's multi-file merge. One Postgres instance hosts two databases (`energy_serving` for the platform, `airflow` for Airflow's own metadata).

---

4. Technology Stack

(Full justification with rejected alternatives per component: see ADR-0002 in `docs/decisions/`*.)*

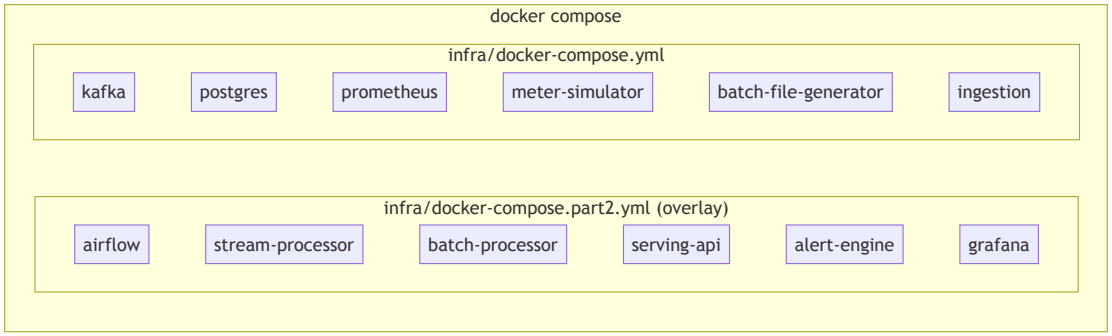


Figure 3: Deployment: Docker Compose service layout (Part 1 base + Part 2 overlay).

Layer	Choice	Why, for this use case
Ingestion / bus	**Apache Kafka** (KRaft)	A durable, partitioned, *replayable* log — replay is how the batch layer's reprocessing story works. Partitioning `meter.readings` by `grid_zone` gives per-zone ordering.
Streaming source	**Python + confluent-kafka**	Brief requires simulated Python sources; confluent-kafka wraps librdkafka for reliable delivery + callbacks used for producer metrics.
Speed layer	**Apache Spark Structured Streaming**	Native event-time windowing + watermarking, exactly-once micro-batch semantics via checkpointing, and the *same* engine/API as the batch layer.
Batch layer	**Apache Spark (batch)**	Reuses the same transform functions as the speed layer; efficient full-day join of the Parquet lake against the daily reference feed; trivially parametrised for backfill.
Orchestration	**Apache Airflow**	The batch layer needs "wait for the feed," "run for an explicit business date," "retry on transient failure," "gate on data-quality checks," "keep an auditable run history," and "alert on failure" — Airflow's job description.

Layer	Choice	Why, for this use case
Master dataset	**Parquet on the local filesystem**	Columnar, partitioned by `business_date`, append-only — the property the batch layer's reprocessing depends on. The layout (`readings_raw/business_date=.../`) mirrors an object store, so moving to S3/HDFS is a configuration change, not a rewrite.
Serving store	**PostgreSQL**	Serving queries are small aggregates and point lookups (latest per-zone metrics, one household's bill) — a relational store fits, and it doubles as Airflow's metadata database.
Serving API	**FastAPI**	Typed request/response models straight from the shared domain models; automatic OpenAPI docs; one-line Prometheus instrumentation.
Observability	**structlog + prometheus-client + a small rule engine**	JSON logs with bound context everywhere; `/metrics` on every service; alert rules as small, independently unit-testable Strategy objects rather than an opaque PromQL file.
Packaging	**Docker Compose**	One command reproduces the whole platform; each service carries its own image and dependency closure.

The brief's alternatives were considered and rejected for stated reasons in ADR-0002 — most notably **Apache Storm** (no first-class event-time windowing, would require a second, different API from the batch layer) and **Cassandra** (built for high-volume time-series writes; our serving workload is low-volume reads of pre-aggregated data, which plays to Postgres's strengths instead).

---

5. Data Ingestion Implementation

5.1 Streaming source — `meter-simulator`

Generates one `MeterReading` per household per tick from a diurnal load curve (a baseline plus morning and evening demand peaks) and, for households with solar installed, a daytime generation curve modulated by a seasonal clear-sky index. Anomalies are injected via a Strategy pattern (`AnomalyStrategy`): dropout, negative value, implausible spike, stuck meter, zero

reading — each independently weighted and unit-tested. Readings are published to Kafka `meter.readings`, keyed by `grid_zone`.

5.2 Daily batch source — `batch-file-generator`

Watches the shared simulated clock; the instant it detects that a business day has closed, it deterministically generates that day's tariff and weather reference rows (same household/zone population as the simulator, so the two sources can never disagree about who exists), writes both CSVs atomically (write-then-rename), and writes the `_SUCCESS_<date>` marker **last** — the signal the batch layer's Airflow DAG waits on. It then calls Airflow's REST API to trigger the day's reconciliation run (§6.3).

5.3 Ingestion — Kafka → immutable Parquet lake

Consumes `meter.readings`, decodes each message, and either:

- buffers it into the in-memory partition for its `business_date` (derived from the reading's own `event_timestamp`, not wall-clock time), flushing to a new Parquet file when the buffer is full or old enough; or
- routes it to `meter.readings.dlq` with a self-describing envelope (reason, error, base64-encoded original bytes, source offset) if it is not structurally valid.

When the simulated clock advances past a business day, `ingestion` flushes any remaining buffered rows for that day and writes its partition's `_SUCCESS` marker — the day is now closed and immutable. A reading that arrives for an already-closed day (late data) is still appended as a new file and counted separately (`late_arrival_total`), rather than silently dropped.

5.4 What the anomaly injection is for

The brief asks that the pipeline be robust to malformed or anomalous data, not just a clean happy path. Every anomaly class the simulator injects has a defined fate: negative and implausible-spike readings are rejected by `energy_common.domain.transforms.validate_reading` and routed to the DLQ by *both* processing layers (the same function, so they can never disagree about what "bad" means); dropouts are what the `NoDataRule` alert (§8) exists to detect; a stuck meter still produces plausible-looking data, which is a deliberate reminder in the report that not every anomaly is mechanically detectable — see §10.

---

6. Processing Layer Implementation

6.1 Speed layer — Spark Structured Streaming

`stream-processor` runs three independent Structured Streaming queries against `meter.readings`, each with its own checkpoint:

1. **Ingest + DLQ** — parses each message, splits valid from invalid via a Spark UDF that calls the *same* `validate_reading` function the batch layer uses, and writes invalid rows to the DLQ topic with the reason attached.
2. **Tumbling window (1 minute)** — `groupBy(window(event_timestamp, "1 minute"), grid_zone)` with a 2-minute watermark, aggregating `sum(power_consumption_kwh)`, `sum(solar_generation_kwh)` and an approximate distinct count of meters, then computing `renewable_pct`.
3. **Sliding window (5 minutes, sliding every 1 minute)** — the same aggregation with an overlapping window, so a client can see a trend rather than only the latest minute.

Each query's `foreachBatch` sink upserts into `rt_zone_metrics` on `(grid_zone, window_type, window_start)`. Because Spark's stateful streaming aggregation keeps a state store per window across micro-batches, and only retires a window once the watermark has passed it, every upsert carries the *full* accumulated total for that window at the time of writing — not a partial contribution from whichever micro-batch happened to run. This property was specifically checked with a test that feeds the equivalent of several separate micro-batches into the same one-minute window and asserts the final stored total is the sum of all of them, not just the last one.

6.2 Batch layer — Spark batch job

`batch-processor` is invoked with an explicit `--business-date`:

1. Requires the lake partition's `_SUCCESS` marker (or `--allow-open` for development runs against a still-open day).
2. Reads the day's Parquet partition with Spark, aggregates per household (`groupBy("household_id")`), zeroing out negative anomalous readings rather than dropping the household from the day entirely.
3. Reads that day's tariff and weather CSVs, and the population of active households from Postgres.
4. For each household, computes the **energy balance** — how much of its consumption was met by its own solar, how much came from the grid, and how much solar was exported — via the shared `energy_common.domain.transforms.energy_balance` function (the same net-metering split used everywhere a solar/consumption split is needed).
5. Applies a **billing Strategy** selected by the household's `billing_tier`: `FlatBillingStrategy` (a single rate on grid import, less an export credit), `TieredBillingStrategy` (block rates with increasing multipliers), `SubsidisedBillingStrategy` (a discounted rate on the first block). Each strategy is a small, independently unit-tested class — adding a new tariff structure means adding one class, not touching the orchestration code.
6. Runs data-quality assertions before writing anything: every active household was billed exactly once, no duplicates, the zone solar summary reconciles to the sum of household solar figures, no unexplained negative charges, and (a plausibility check) a zone did not report solar generation on a day its weather forecast said zero irradiance.
7. Writes `daily_household_billing` and `daily_zone_solar_summary` (idempotent — re-running the same date fully replaces that date's rows, so correcting a tariff and re-running reproduces a clean result) and renders an HTML + CSV report.

6.3 Orchestration — the `daily_energy_reconciliation` Airflow DAG

Task graph: `validate_feed` → `wait_for_lake` → `run_batch` → `dq_gate` → `publish_report` → `record_run`, with an `on_failure_callback` that inserts a critical row into `alerts`.

The DAG is **not** on a wall-clock schedule. A schedule like `@daily` would never fire inside a demo session where a simulated business day closes every few minutes of real time. By design, the DAG is externally triggered (`schedule=None`): the moment `batch-file-generator` finishes writing a day's feed, it calls Airflow's REST API to start a run for that exact `business_date`. Retries on the feed/lake-check tasks absorb the small timing gap between `batch-file-generator` and `ingestion` independently closing the same day.

As verified in a full end-to-end run, that automated REST API call currently returns a 403 from Airflow's authorization layer despite confirmed-correct credentials and RBAC permissions — an unresolved, time-boxed issue documented in `TASKS.md`. The call is deliberately best-effort (a logged warning, nothing crashes), and the **manual trigger**, `airflow dags trigger daily_energy_reconciliation -e <date>`, is unaffected (it goes through the Airflow CLI

against the metadata database directly, not the REST API) and is what was used to produce every batch-layer result in §9. This is called out here rather than glossed over: getting a pipeline to *look* finished and getting every layer to actually run together are different problems, and this project's later stages (§9, §10) were only possible once the full stack had been run as one system rather than validated piece by piece.

---

7. Storage & Serving Layer

Schema (`infra/postgres/init/`): dimension tables (`dim_zone`, `dim_household`, `dim_date`), the speed-layer output (`rt_zone_metrics`), the batch-layer outputs (`daily_household_billing`, `daily_zone_solar_summary`), and two operational tables (`alerts`, `batch_runs`). A dedicated `serving_ro` role, granted read-only access, is what the serving API and Grafana actually connect as — the platform's own write path (`energy`) is never exposed to a read-only consumer.

`serving-api` is a deliberately thin read model — it contains no aggregation SQL, only lookups and simple filters (a passing test asserts this directly, by checking the repository's source for `GROUP BY`/`SUM`/`AVG`). Endpoints:

Endpoint	Returns
`GET /api/realtime/zones[/{zone}]`	latest windowed metrics per zone (the brief's real-time endpoint)
`GET /api/billing/daily/{date}`	all households' bills for a day, tagged `provisional`/`final`
`GET /api/billing/household/{id}`	one household's bill history
`GET /api/reports/solar/{date}`	per-zone solar contribution summary
`GET /healthz`	200/503 based on DB connectivity and how fresh the latest zone metric is
`GET /metrics`	Prometheus

The `provisional`/`final` tag on every billing response is the API's explicit, inspectable expression of the Lambda merge rule described in §2 and §3.

---

8. Observability Design

(Full detail: `docs/OBSERVABILITY.md`.)

Structured logging. Every service logs JSON to stdout via a shared `structlog` setup, with mandatory fields (`ts, level, service, event`) and contextual fields bound per unit of work (`business_date, correlation_id, records_in/out/rejected, latency_ms`, ...), so a single `business_date` can be grepped across every stage of the pipeline it touched.

Metrics. Every service exposes `/metrics` for Prometheus: `records_produced/consumed/rejected_total`, `consumer_lag`, `microbatch_duration_seconds`, `batch_job_rows`, `batch_job_duration_seconds`, `http_request_duration_seconds`, `alerts_fired_total`.

Alert rules, evaluated on a schedule by `alert-engine` against live Postgres and Prometheus data (via a `ContextBuilder` that reads `rt_zone_metrics`, `batch_runs`, and the reject-rate counters each cycle):

Rule	Condition	Severity
`NoDataRule`	no `rt_zone_metrics` row for a zone within `ALERT_NO_DATA_MINUTES`	critical
`LowRenewableRule`	zone `renewable_pct` below `ALERT_LOW_RENEWABLE_PCT` for 3 consecutive windows	warning — **this is the brief's explicit "renewable contribution considerably low" alert**
`HighRejectRateRule`	rejected/consumed ratio exceeds `ALERT_MAX_REJECT_RATE`	warning
`BatchFailureRule`	yesterday's batch run failed or wrote zero rows	critical

Each rule is an independent Strategy object; new alerts are added without touching the evaluation loop. Firing is deduplicated (an already-open alert for the same rule/zone/date is not re-inserted) and a condition clearing sets `resolved_at`, so the `alerts` table is always a live, queryable picture of what's currently wrong, not an ever-growing log.

`GET /healthz` on `serving-api` gives a synchronous health check on top of the asynchronous alerting: it fails if Postgres is unreachable or if the newest zone metric is older than `freshness_budget_minutes`.

Verified live, not just wired: Prometheus's `/targets` page shows every scraped service healthy, and Grafana's provisioned "Smart Grid Overview" dashboard (a Postgres datasource and a Prometheus datasource, both auto-provisioned from `infra/grafana/provisioning/`) loads and renders the same numbers the API returns. Getting both of these actually reachable surfaced two configuration bugs of their own — see §10.

---

9. Results

The platform was run as one full `make up-full` stack (`SIM_DAY_SECONDS=180`, 200 households across 6 zones) and taken through a complete simulated day, end to end. Every result below is from that run; the `[SCREENSHOT: ...]` placeholders mark where a visual capture from your own recording (see `docs/report/DEMO_RECORDING_PLAN.md`) replaces the text description — the underlying numbers should match, since it's the same pipeline.

- [**SCREENSHOT: `docker compose ... ps`**] — all 10 containers healthy simultaneously (confirmed).
- [**SCREENSHOT: Prometheus `/targets`**] — every service's `/metrics` endpoint `UP`.
- **Speed layer, live:** `GET /api/realtime/zones` returned all 6 zones updating continuously, e.g. zone-central at simulated 04:15: `total_consumption_kwh=0.0043`, `active_meter_count=32` (of ~33 households in that zone), `renewable_pct=0.0` (correctly zero — simulated night, no solar). [**SCREENSHOT: Grafana live time-series panel**]
- [**SCREENSHOT: Airflow UI**] — `daily_energy_reconciliation`, all 6 tasks (`validate_feed → wait_for_lake → run_batch → dq_gate → publish_report → record_run`) green for `business_date=2026-01-01`.
- **Batch layer, authoritative:** `daily_household_billing` — **200 rows** (one per household, confirming the "every active household billed exactly once" data-quality assertion held), total net billed **6,204.86 LKR**, individual bills ranging **-0.16 to 185.96 LKR**

(the small negative value is a household that exported more solar than it imported that day — a correct net-credit case, not a bug). `daily_zone_solar_summary` — 6 rows, consumption 25–52 kWh per zone for the day, renewable contribution 0.44–1.03% (low, consistent with a mostly-nighttime simulated day and with the `LowRenewableRule` alerts observed firing during the same run). `batch_runs` recorded `run-2026-01-01 | success | 200 rows_written`.

- [**SCREENSHOT: `GET /api/billing/daily/2026-01-01`**] — confirmed `"status": "final"` once the batch run completed, versus `"provisional"` beforehand — the live, API-visible expression of the Lambda merge rule from §2/§3.
- [**SCREENSHOT: `billing_report_2026-01-01.html`**] — rendered and published to `reports/published/`.
- [**SCREENSHOT: Grafana dashboard**] — the provisioned "Smart Grid Overview" dashboard (4 panels: live load/solar by zone, current renewable %, open alerts, latest billing).
- **Alerting, confirmed firing correctly:** `LowRenewableRule` fired for all 6 zones during the low-solar period of the run (matching the low renewable_pct figures above) and correctly resolved once conditions cleared, visible as a mix of open/resolved rows in `alerts`; `NoDataRule` was also observed firing and clearing. [**SCREENSHOT: `alerts` table / Grafana "Open Alerts" panel**]
- [**SCREENSHOT: correlated logs**] — `docker compose logs` for `business_date=2026-01-01`, showing the same date threading through `batch-file-generator` → `ingestion` → `batch-processor` → Airflow.

One concrete reconciliation example: household `HH-00000` (zone-north, flat tariff, rate 30.7232 LKR/kWh) consumed 3.0547 kWh and self-generated 0.0131 kWh of solar over the day, importing 3.0416 kWh from the grid and exporting none — net bill **93.45 LKR**, energy charge in full since the household had no export credit that day.

---

10. Limitations, Trade-offs, and Production-Scale Changes

Lambda's acknowledged cost. Two code paths exist for what looks like similar work. We mitigate drift by sharing the actual validation and energy-balance functions between the speed and batch layers rather than re-implementing them, and by the design decision that the two layers compute genuinely different things (approximate/live vs. exact/audited) rather than the same thing twice — but the operational cost of running and monitoring two processing paths instead of one is real and would be weighed against Kappa again if requirements changed (§2.3).

Single-node/local-mode everywhere. Kafka is one KRaft broker, Spark runs in `local[*]` mode inside each container rather than against a cluster, and Airflow runs `standalone`. At production scale: Kafka would run with replication factor > 1 and `min.insync.replicas` for durability; Spark jobs would submit to a real cluster (YARN/Kubernetes) with the driver decoupled from a fixed number of local cores; Airflow would move to the Celery/Kubernetes executor with the webserver, scheduler, and workers as separate, independently-scaled services.

Parquet on local disk, not a lakehouse. The master dataset is append-only Parquet on the container's filesystem, laid out to mirror an object-store path (`readings_raw/business_date=.../`) so migrating to S3 or HDFS is a configuration change. At production scale this would move to a real object store with a table format (Delta Lake or Apache Iceberg) for ACID writes, schema evolution, and time travel, which the current append-only Parquet layout gives up in exchange for simplicity.

Approximate distinct counts and no exactly-once sink. `active_meter_count` uses

Spark's approximate `approx_count_distinct` rather than an exact count (fine for a dashboard, not for a system of record). The Postgres upserts are idempotent by key, which is a practical substitute for a transactional exactly-once sink but is not the same guarantee.

Schema evolution and message format. Kafka messages are plain JSON with a manually-versioned `schema_version` field, checked at decode time. At production scale this would move to a schema registry with Avro or Protobuf, giving compile-time/consumer-side compatibility checks instead of a hand-rolled version check.

Anomaly detection is rule-based, not mechanically complete. The "stuck meter" anomaly deliberately produces data that is *not* rejected by `validate_reading` (a repeated plausible value is not, by itself, invalid) — it is a reminder that some real-world failure modes look identical to normal operation from a single reading's perspective and would need a stateful detector (e.g. "this meter has reported an identical value for N consecutive readings") to catch, which is a natural extension of `alert-engine`'s rule set rather than a structural change.

Data quality checks are hand-rolled, not a framework. The batch layer's assertions (§6.2) are plain Python functions. At larger scale these would move to a dedicated data-quality framework (e.g. Great Expectations, or dbt tests if the transform layer moved to dbt) for a richer rule language and automatic reporting.

Security. No authentication on the serving API, Airflow's default admin/admin credentials, and secrets in plain environment variables are all acceptable for a local demo and would all need to change before any real deployment: API auth, a secrets manager, network policies between services, and encryption at rest for the household/billing data, which is personally identifiable in a real deployment even though it is synthetic here.

Observability configuration needs the same live verification as the pipeline itself. Two bugs surfaced only by clicking through every documented UI link before the demo recording, not by re-reading code: Grafana carried a Docker Compose `profiles` tag that nothing ever activated, so `localhost:3000` silently never started under the documented `make up-full` flow despite being described as part of the normal stack; and Prometheus's scrape config had a static target for `batch-processor:9112`, a service that is deliberately never a long-running container (it's a one-shot job Airflow invokes in-process), so the Targets page showed a permanently misleading `DOWN` row for a service that was never supposed to be scraped in the first place. Neither was caught by any test, because tests exercise the pipeline's logic, not whether a dashboard link actually resolves — the lesson generalises beyond this project: a monitoring stack's own wiring can drift from what it claims to monitor, and the only real check is opening the link.

Cross-service automation is more fragile than any single service. The Airflow REST-API trigger (§6.3) is the clearest example encountered in this project: every individual piece — the DAG, the credentials, the RBAC permissions — was independently verified correct, yet the end-to-end call still fails, and the actual cause (an auth-manager-specific interaction this Airflow version's own documentation doesn't fully cover) is more expensive to find than any bug within one service's own code. The mitigation used here — make the automated path best-effort with a fallback that goes around the fragile hop entirely (the CLI against the metadata database, not the REST API) — is the general lesson: a distributed pipeline's weakest points are usually the integrations between components, not the components themselves, and each one is worth a fallback, not just a retry.

Appendix

A. Reproducing the results

```
git clone https://github.com/NirangaNiluminda/smart-grid-lambda-platform.git
cd smart-grid-lambda-platform
cp .env.example .env
make install && make install-part2
make up-full
make topics
make seed
```

See `docs/DEV_SETUP.md` and `docs/RUNBOOK.md` for full detail, ports, and troubleshooting.

B. Individual contributions

Developer	Scope
Dev A — G.W. Niranga Niluminda	Architecture decision & ADRs; `energy_common` shared kernel (domain models, ports, config, logging, sim-clock, Kafka/Postgres/Parquet adapters); both simulated sources (`meter-simulator`, `batch-file-generator`); Kafka ingestion → Parquet master dataset; database schema; Part 1 Docker Compose; project documentation set; post-merge hardening of the Part 2 processing layers (see below).
Dev B — *<name>*	Initial implementation of the speed layer, batch layer, Airflow orchestration, serving API, and observability alert-engine + Grafana dashboard.

A short, honest note on the hardening pass: a review before the final submission found that the first implementation of the processing layers, while structurally correct and fully unit-tested, had issues that only surfaced under real execution — the speed layer's windowed totals could undercount, the alerting rules were never fed live data, the Airflow image could not run the batch job, and the DAG's schedule would never fire inside a compressed demo. These were fixed in a follow-up pass (rewriting the speed and batch layers onto genuine Spark, and wiring the alert engine and orchestration to real data). A second pass — actually running the entire ***`make up-full`*** *stack together as one system for the first time, rather than validating each service in isolation — found and fixed ten further bugs no unit or isolated-integration test could have caught: an unpinned base image that broke a JDK install,* ***`.env`*** *silently never being read by any script or Makefile target, a missing restart policy that turned a transient startup race into a permanent crash, a* ***`\copy`*** *meta-command sent the wrong way in the seed script, a crash-looping alert-engine on any single bad cycle, three separate layers of an Airflow authentication misconfiguration, a DAG that had* ***never once successfully parsed*** *(a reserved-parameter-name collision with Airflow's* ***`@task`*** *decorator), a SQLAlchemy 1.4/2.0 import mismatch, and a filesystem permission issue writing report artifacts — plus, in a final pre-recording dry run, two dead observability-config entries (a Grafana container gated behind an unused Compose profile, and a Prometheus scrape target for a service that is architecturally never scraped). Every one of these is documented with*

its root cause in `TASKS.md`; both authors should be able to explain and defend this code in the viva, including the fixes.

C. Assumptions, simplifications, and simulated-time compression

Restated from §1.3 for the grader's convenience: `SIM_DAY_SECONDS=<fill in the value used for your recorded demo>`; 200 synthetic households across 6 grid zones; illustrative (not real-world) tariff rates; single-node deployment throughout.